

# Multi Linear Regression

## Polynomial Regression

### Linear Regression with Multiple Variables

- input = more than 1 features
- output = single target
- get or load the data
- data preprocessing
- define input and output
- apply the model or algorithm
- apply train/test data
- evaluate the score
- $y = ax^2 + bx + c$  degree=2
- $y = ax^3 + bx^2 + cx + 1$  degree=3

```
In [1]: 1 # importing libraries
        2 import numpy as np
        3 import pandas as pd
        4 import matplotlib.pyplot as plt
        5
```

### Prediction of the house price of boston dataset

#### 1.get the data

```
In [2]: 1 from sklearn.datasets import load_boston
```

```
In [3]: 1 # create object
        2 boston = load_boston()
        3
```

```
In [4]: 1 boston.keys()
```

```
Out[4]: dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
```

```
In [5]: 1 print(boston['DESCR'])
```

```
.. _boston_dataset:
```

```
Boston house prices dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
    :Number of Instances: 506
```

```
    :Number of Attributes: 13 numeric/categorical predictive. Median Value (a
    ttribute 14) is usually the target.
```

```
    :Attribute Information (in order):
```

```
        - CRIM      per capita crime rate by town
        - ZN        proportion of residential land zoned for lots over 25,000
sq.ft.
        - INDUS     proportion of non-retail business acres per town
        - CHAS      Charles River dummy variable (= 1 if tract bounds river; 0
otherwise)
        - NOX       nitric oxides concentration (parts per 10 million)
        - RM        average number of rooms per dwelling
        - AGE       proportion of owner-occupied units built prior to 1940
        - DIS       weighted distances to five Boston employment centres
        - RAD       index of accessibility to radial highways
        - TAX       full-value property-tax rate per $10,000
        - PTRATIO   pupil-teacher ratio by town
        - B         1000(Bk - 0.63)^2 where Bk is the proportion of blacks by
town
        - LSTAT     % lower status of the population
        - MEDV      Median value of owner-occupied homes in $1000's
```

```
    :Missing Attribute Values: None
```

```
    :Creator: Harrison, D. and Rubinfeld, D.L.
```

```
This is a copy of UCI ML housing dataset.
```

```
https://archive.ics.uci.edu/ml/machine-learning-databases/housing/ (https://a
rchive.ics.uci.edu/ml/machine-learning-databases/housing/)
```

```
This dataset was taken from the StatLib library which is maintained at Carneg
ie Mellon University.
```

```
The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic
prices and the demand for clean air', J. Environ. Economics & Management,
vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics
...', Wiley, 1980. N.B. Various transformations are used in the table on
pages 244-261 of the latter.
```

```
The Boston house-price data has been used in many machine learning papers tha
t address regression
problems.
```

```
.. topic:: References
```

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [7]: 1 `len(boston['feature_names'])`

Out[7]: 13

In [8]: 1 `boston['data']`

Out[8]: array([[6.3200e-03, 1.8000e+01, 2.3100e+00, ..., 1.5300e+01, 3.9690e+02, 4.9800e+00],  
[2.7310e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9690e+02, 9.1400e+00],  
[2.7290e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9283e+02, 4.0300e+00],  
...,  
[6.0760e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02, 5.6400e+00],  
[1.0959e-01, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9345e+02, 6.4800e+00],  
[4.7410e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02, 7.8800e+00]])

In [9]: 1 `df = pd.DataFrame(boston['data'])`

In [10]: 1 `df`

Out[10]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33
...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.67
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.08
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.64
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.48
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.88

506 rows × 13 columns

```
In [11]: 1 df.columns = boston['feature_names']
```

```
In [12]: 1 df
```

```
Out[12]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.
...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.

506 rows × 13 columns



```
In [13]: 1 df['target']=boston['target']
```

```
In [14]: 1 df
```

```
Out[14]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST	target
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.	5.05
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.	16.99
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.	16.99
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.	16.99
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.	16.99
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.	16.99
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.	16.99
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.	16.99
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.	16.99
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.	16.99

506 rows × 14 columns



In [15]: 1 df.shape

Out[15]: (506, 14)

In [16]: 1 # 2.Preprocessing the data  
2 df.isna().sum()

Out[16]: CRIM 0  
ZN 0  
INDUS 0  
CHAS 0  
NOX 0  
RM 0  
AGE 0  
DIS 0  
RAD 0  
TAX 0  
PTRATIO 0  
B 0  
LSTAT 0  
target 0  
dtype: int64

In [18]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        506 non-null    float64
1   ZN          506 non-null    float64
2   INDUS       506 non-null    float64
3   CHAS        506 non-null    float64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         506 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    float64
9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
13  target      506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB
```

In [19]: 1 df.describe()

Out[19]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.00
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.79
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.10
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.12
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.10
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.20
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.18
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.12

In [20]:

```

1 # 3. Define input and output
2 X = df[['RM']]
3 y = df['target']

```

- it is better to separate the data for training data and testing data
- we can say 70% data for training and 30% data for testing
- we have 506 rows available in dataframe
- how many rows for training and how many rows for testing

In [21]: 1 70\*506/100 # training data

Out[21]: 354.2

In [22]:

```

1 506-354 # testing data
2

```

Out[22]: 152

In [23]: 1 from sklearn.model\_selection import train\_test\_split

In [24]: 1 X\_train,X\_test,y\_train,y\_test = train\_test\_split(X,y,train\_size=0.7)

In [26]: 1 X\_train.shape

Out[26]: (354, 1)

In [27]: 1 X\_test.shape

Out[27]: (152, 1)

```
In [28]: 1 y_train.shape
```

```
Out[28]: (354,)
```

```
In [29]: 1 y_test.shape
```

```
Out[29]: (152,)
```

```
In [30]: 1 # train the model
2 from sklearn.linear_model import LinearRegression
```

```
In [31]: 1 model = LinearRegression()
```

```
In [32]: 1 model.fit(X_train,y_train)
```

```
Out[32]: LinearRegression()
```

```
In [33]: 1 print('training score',model.score(X_train,y_train)*100)
```

```
training score 45.73896073809952
```

```
In [34]: 1 print('testing score',model.score(X_test,y_test)*100)
```

```
testing score 54.340073017363125
```

improve the model above score very low so want to improve the model 1.giving the more examples 2.by taking different features 3.by parameter tuning

In [35]: 1 df.corr()

Out[35]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	-0.379670	0.
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	0.664408	-0.
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	-0.708027	0.
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	-0.099176	-0.
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	-0.769230	0
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	0.205246	-0.
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	-0.747881	0.
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	1.000000	-0.
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	-0.494588	1.
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	-0.534432	0.
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	-0.232471	0.
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	0.291512	-0.
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	-0.496996	0.
target	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	0.249929	-0.

In [50]: 1 X = df[['RM', 'PTRATIO', 'LSTAT']]  
2 y = df['target']

In [51]: 1 from sklearn.model\_selection import train\_test\_split

In [53]: 1 X\_train, X\_test, y\_train, y\_test = train\_test\_split(X, y, train\_size=0.7)

In [54]: 1 X\_train.shape

Out[54]: (354, 3)

In [55]: 1 from sklearn.linear\_model import LinearRegression  
2 model = LinearRegression()  
3 model.fit(X\_train, y\_train)

Out[55]: LinearRegression()

In [56]: 1 model.score(X\_train, y\_train)\*100

Out[56]: 71.1808054134889



```
In [58]: 1 model.score(X_test,y_test)*100
```

```
Out[58]: 59.86843710824545
```

```
In [67]: 1 X=df.drop('target',axis=1)
```

```
In [68]: 1 y=df['target']
```

```
In [69]: 1 from sklearn.model_selection import train_test_split
```

```
In [70]: 1 X_train,X_test,y_train,y_test = train_test_split(X,y,train_size=0.7)
```

```
In [71]: 1 X_train.shape
```

```
Out[71]: (354, 13)
```

```
In [72]: 1 from sklearn.linear_model import LinearRegression  
2 model = LinearRegression()  
3 model.fit(X_train,y_train)
```

```
Out[72]: LinearRegression()
```

```
In [73]: 1 model.score(X_train,y_train)*100
```

```
Out[73]: 74.53585738547092
```

```
In [85]: 1 # Applying polynomial features to linear regression  
2  
3 experience = [0,1,2,3,4,5,6,7,8]  
4 salary = [5000,6000,7000,8000,15000,25000,40000,50000,80000]
```

```
In [86]: 1 df = pd.DataFrame({'experience':experience,'salary':salary})
```

In [87]:

```
1 df
```

Out[87]:

	experience	salary
0	0	5000
1	1	6000
2	2	7000
3	3	8000
4	4	15000
5	5	25000
6	6	40000
7	7	50000
8	8	80000

In [88]:

```
1 df.shape
```

Out[88]: (9, 2)

In [89]:

```
1 df.isna().sum()
```

Out[89]:

experience	0
salary	0

dtype: int64

In [90]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9 entries, 0 to 8
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   experience  9 non-null      int64
1   salary      9 non-null      int64
dtypes: int64(2)
memory usage: 272.0 bytes
```

In [91]: 1 df.describe()

Out[91]:

	experience	salary
count	9.000000	9.000000
mean	4.000000	26222.222222
std	2.738613	25825.267558
min	0.000000	5000.000000
25%	2.000000	7000.000000
50%	4.000000	15000.000000
75%	6.000000	40000.000000
max	8.000000	80000.000000

In [92]: 1 X=df[['experience']]  
2 y = df['salary']

In [93]: 1 X\_train =X.head(7)  
2 X\_test =X.tail(2)  
3 y\_train = y.head(7)  
4 y\_test = y.tail(2)


In [94]: 1 from sklearn.linear\_model import LinearRegression  
2 model = LinearRegression()  
3 model.fit(X\_train,y\_train)

Out[94]: LinearRegression()

In [95]: 1 model.score(X\_train,y\_train)\*100

Out[95]: 79.92498597868762

In [97]: 1 plt.scatter(df['experience'],df['salary'],c='blue',label = 'true values')  
2 plt.plot(df['experience'],model.predict(X),c='red',label = 'predicted line')  
3 plt.show()



```
In [98]: 1 from sklearn.preprocessing import PolynomialFeatures
```

```
In [99]: 1 poly = PolynomialFeatures()  
2
```

```
In [100]: 1 X_poly_train = poly.fit_transform(X_train)
```

```
In [101]: 1 X_poly_test = poly.transform(X_test)
```

```
In [102]: 1 from sklearn.linear_model import LinearRegression  
2
```

```
In [103]: 1 model = LinearRegression()
```


```
In [104]: 1 model.fit(X_poly_train,y_train)
```

Out[104]: LinearRegression()

```
In [105]: 1 model.score(X_poly_train,y_train)*100
```

Out[105]: 98.77079828005235

```
In [106]: 1 plt.scatter(df['experience'],df['salary'],c='green',label='true values')  
2 plt.plot(df['experience'],model.predict(poly.transform(X)),c='red',label='pr  
3 plt.show()
```



```
In [ ]: 1
```