

Today Concepts

- Polynomial Regression with Multiple Features
- Ridge Regression
- Lasso Regression

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
```

```
In [2]: 1 df = pd.read_csv("https://raw.githubusercontent.com/AP-State-Skill-Developme
```

```
In [3]: 1 df
```

```
Out[3]:
```

	MODELYEAR	MAKE	MODEL	VEHICLECLASS	ENGINESIZE	CYLINDERS	TRANSMISSION
0	2014	ACURA	ILX	COMPACT	2.0	4	AS5
1	2014	ACURA	ILX	COMPACT	2.4	4	M6
2	2014	ACURA	ILX HYBRID	COMPACT	1.5	4	AV7
3	2014	ACURA	MDX 4WD	SUV - SMALL	3.5	6	AS6
4	2014	ACURA	RDX AWD	SUV - SMALL	3.5	6	AS6
...
1062	2014	VOLVO	XC60 AWD	SUV - SMALL	3.0	6	AS6
1063	2014	VOLVO	XC60 AWD	SUV - SMALL	3.2	6	AS6
1064	2014	VOLVO	XC70 AWD	SUV - SMALL	3.0	6	AS6
1065	2014	VOLVO	XC70 AWD	SUV - SMALL	3.2	6	AS6
1066	2014	VOLVO	XC90 AWD	SUV - STANDARD	3.2	6	AS6

1067 rows × 13 columns

```
In [4]: 1 df.columns
```

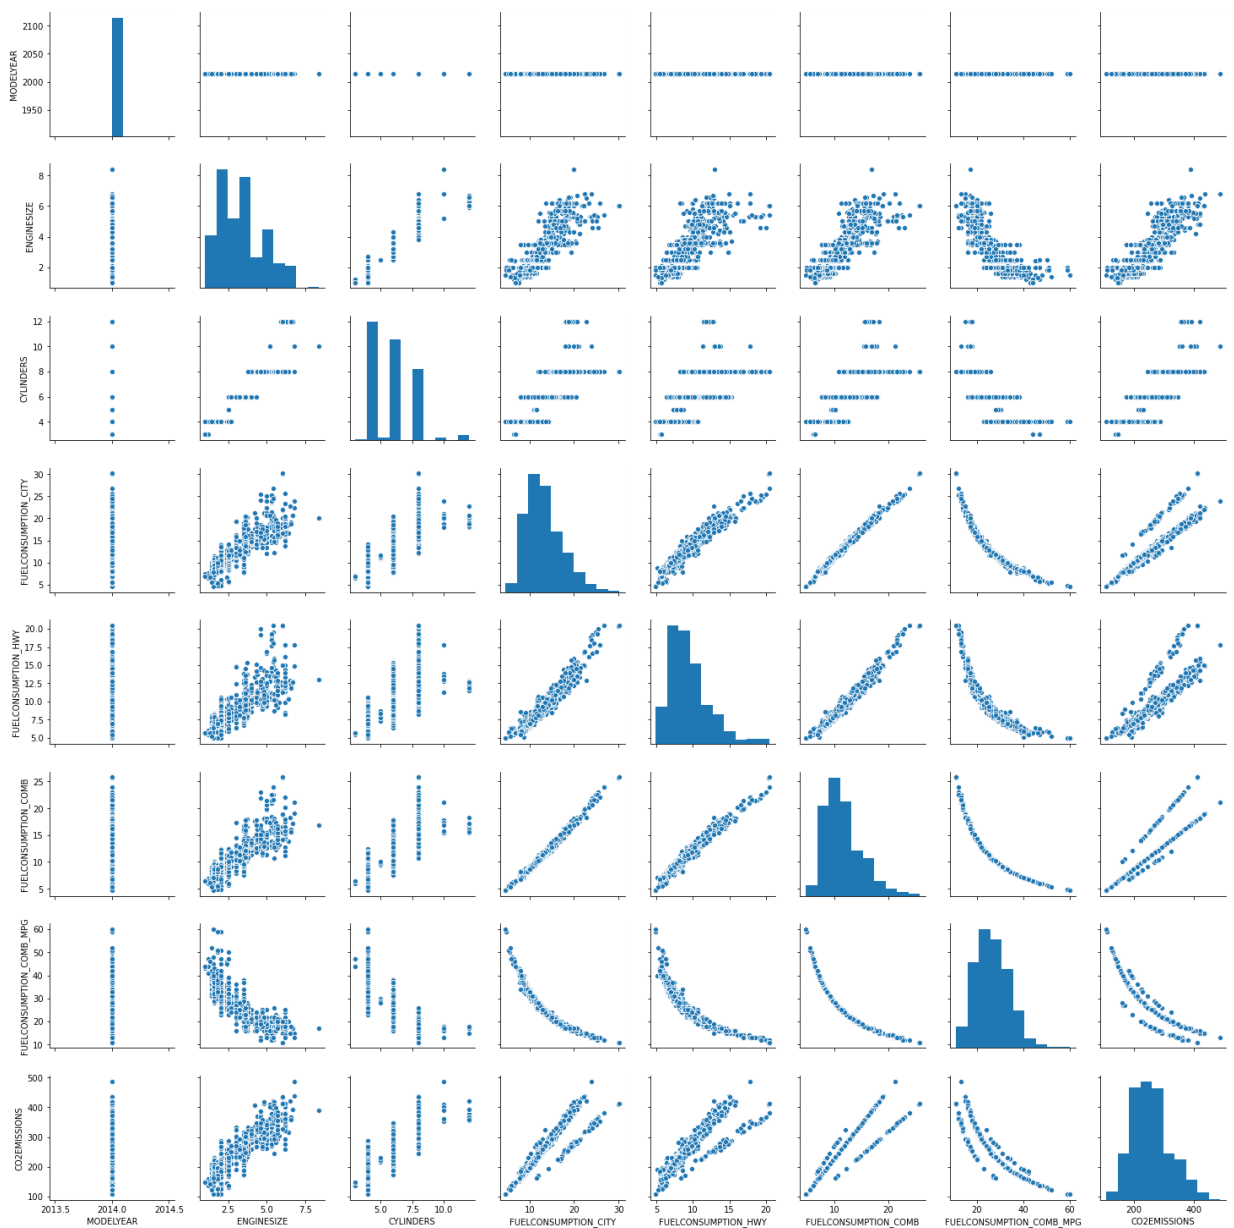
```
Out[4]: Index(['MODELYEAR', 'MAKE', 'MODEL', 'VEHICLECLASS', 'ENGINESIZE', 'CYLINDERS',
               'TRANSMISSION', 'FUELTYPE', 'FUELCONSUMPTION_CITY',
               'FUELCONSUMPTION_HWY', 'FUELCONSUMPTION_COMB',
               'FUELCONSUMPTION_COMB_MPG', 'CO2EMISSIONS'],
              dtype='object')
```

```
In [5]: 1 df['CO2EMISSIONS'].head(10)
```

```
Out[5]: 0    196
1    221
2    136
3    255
4    244
5    230
6    232
7    255
8    267
9    212
Name: CO2EMISSIONS, dtype: int64
```


```
In [6]: 1 import seaborn as sns
2 sns.pairplot(df)
```

```
Out[6]: <seaborn.axisgrid.PairGrid at 0x1f7b66dd1c0>
```



```
In [7]: 1 x = np.arange(-5,5,0.1)
        2 plt.plot(x,x**2,c='r')
```

Out[7]: [<matplotlib.lines.Line2D at 0x1f7bd7ffdc0>]



```
In [21]: 1 X = df[['FUELCONSUMPTION_COMB_MPG']].values.reshape(-1,1)
```

```
In [22]: 1 X
```

Out[22]: array([[33],
[29],
[48],
...,
[24],
[25],
[22]], dtype=int64)

```
In [23]: 1 y = df['CO2EMISSIONS']
```

```
In [24]: 1 from sklearn.preprocessing import PolynomialFeatures
```

```
In [25]: 1 poly = PolynomialFeatures(degree=2)
```

```
In [26]: 1 X_poly = poly.fit_transform(X)
```

```
In [27]: 1 X_poly
```

Out[27]: array([[1.000e+00, 3.300e+01, 1.089e+03],
[1.000e+00, 2.900e+01, 8.410e+02],
[1.000e+00, 4.800e+01, 2.304e+03],
...,
[1.000e+00, 2.400e+01, 5.760e+02],
[1.000e+00, 2.500e+01, 6.250e+02],
[1.000e+00, 2.200e+01, 4.840e+02]])

```
In [28]: 1 from sklearn.linear_model import LinearRegression
```

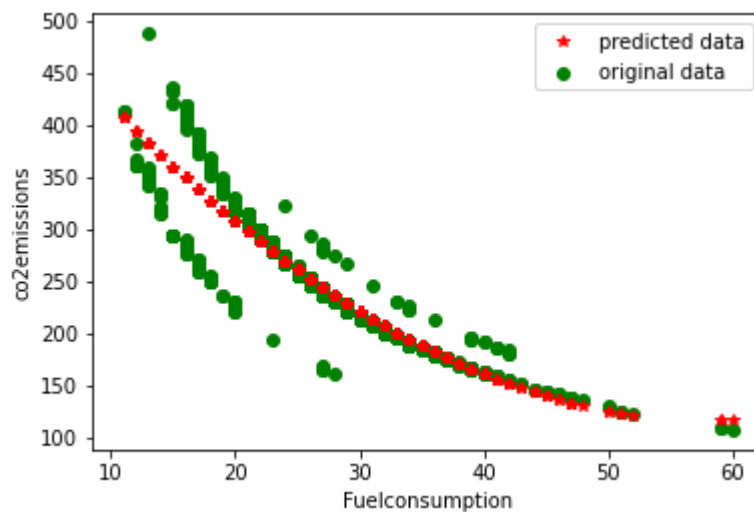
```
In [29]: 1 model = LinearRegression()
```

```
In [30]: 1 model.fit(X_poly,y)
```

```
Out[30]: LinearRegression()
```

```
In [31]: 1 y_pred = model.predict(X_poly)
```

```
In [33]: 1 plt.scatter(X,y,label = 'original data',c='g')
2 plt.plot(X,y_pred,'r*',label = 'predicted data')
3 plt.xlabel('Fuelconsumption')
4 plt.ylabel('co2emissions')
5 plt.legend()
6 plt.show()
```



Ridge Regression

```
In [34]: 1 m=100
2 x = np.random.rand(m,1)
3 y = 0.5*x**2+x+2+np.random.randn(m,1)
```

In [35]: 1 print(x)

```
[0.1275804 ]  
[0.56379826]  
[0.92273745]  
[0.63536456]  
[0.52386528]  
[0.10697579]  
[0.55070548]  
[0.73659658]  
[0.44002359]  
[0.45982298]  
[0.79428537]  
[0.9116974 ]  
[0.5458613 ]  
[0.59877542]  
[0.5374324 ]  
[0.43890212]  
[0.93233761]  
[0.47384011]  
[0.59547886]  
[0.8656434 ]  
[0.87900426]  
[0.69395659]  
[0.1980226 ]  
[0.22291654]  
[0.40345386]  
[0.35787618]  
[0.96015506]  
[0.81998997]  
[0.83728784]  
[0.38720258]  
[0.32799929]  
[0.84188215]  
[0.55821378]  
[0.8810462 ]  
[0.47645051]  
[0.07594555]  
[0.95640042]  
[0.39512064]  
[0.68693279]  
[0.63075027]  
[0.93633961]  
[0.98952706]  
[0.00553965]  
[0.30891415]  
[0.30992322]  
[0.06333435]  
[0.81849027]  
[0.3630223 ]  
[0.58041542]  
[0.13714473]  
[0.36788261]  
[0.14730769]  
[0.86602458]  
[0.36858089]
```

[0.6020984]
[0.72639997]
[0.05977751]
[0.41982494]
[0.40627689]
[0.49813486]
[0.98589343]
[0.88899879]
[0.99063333]
[0.18343994]
[0.30986029]
[0.9045653]
[0.90431475]
[0.23278418]
[0.66641164]
[0.51408005]
[0.27970614]
[0.96444049]
[0.76442354]
[0.39315697]
[0.00757536]
[0.38582148]
[0.92373111]
[0.68696396]
[0.67252587]
[0.07709121]
[0.37533877]
[0.82975991]
[0.89445485]
[0.08411864]
[0.58147322]
[0.57851751]
[0.98477089]
[0.36497697]
[0.7119468]
[0.64416844]
[0.76898914]
[0.66601207]
[0.16657521]
[0.33766315]
[0.42319038]
[0.30438249]
[0.11250165]
[0.8090909]
[0.95969813]
[0.62076427]

In [36]:

1 print(y)

```
[1.77164985]
[2.61666142]
[4.88594726]
[3.79319673]
[0.71991947]
[2.30618844]
[3.89221617]
[2.85840739]
[3.24587906]
[2.69005347]
[4.13484452]
[3.70631309]
[2.81297893]
[2.93964509]
[2.63462184]
[1.58566038]
[2.66781196]
[2.41157247]
[2.29899402]
[3.28415413]
[0.51146693]
[3.0197024 ]
[4.40095372]
[2.47996699]
[2.47611495]
[2.06161329]
[3.43248565]
[2.7108503 ]
[4.24002009]
[2.10475234]
[2.99275358]
[2.02793172]
[1.33056952]
[1.54219218]
[3.52025314]
[2.82511037]
[4.73017684]
[3.1844743 ]
[2.55899585]
[2.42832857]
[1.978807 ]
[3.30621499]
[1.13086522]
[3.22213357]
[1.95825311]
[2.60019126]
[3.15595027]
[1.53692436]
[2.84056513]
[1.78002929]
[1.51350938]
[2.47172943]
[0.55894746]
[1.99632419]
```

```
[4.21242025]  
[2.14576266]  
[3.94756345]  
[2.83364223]  
[2.01844412]  
[1.75315918]  
[3.64138443]  
[2.57916464]  
[4.37832628]  
[0.56120929]  
[5.17711629]  
[2.44143702]  
[3.46837417]  
[2.47761714]  
[2.68283687]  
[2.89336637]  
[3.30306414]  
[3.08511995]  
[3.47874408]  
[1.43744338]  
[1.69219673]  
[0.61281211]  
[2.56663093]  
[2.93496943]  
[2.57310399]  
[2.06096387]  
[2.86520124]  
[4.01524309]  
[1.86249689]  
[1.37929831]  
[2.64952666]  
[1.39850018]  
[2.80246834]  
[1.46213683]  
[3.57483405]  
[2.71251131]  
[1.51342583]  
[2.69844427]  
[0.72210201]  
[2.43477925]  
[1.96750145]  
[3.37074469]  
[3.76013114]  
[3.18557085]  
[2.42767396]  
[3.01974994]]
```

In [37]: 1 x.shape

Out[37]: (100, 1)

In [38]: 1 y.shape

Out[38]: (100, 1)


```
In [39]: 1 from sklearn.linear_model import Ridge
        2 rd = Ridge()
```

```
In [40]: 1 x[1]
```

```
Out[40]: array([0.56379826])
```

```
In [41]: 1 np.ndim(x)
```

```
Out[41]: 2
```

```
In [42]: 1 rd.fit(x,y)
```

```
Out[42]: Ridge()
```

```
In [44]: 1 rd.score(x,y)*100
```

```
Out[44]: 6.989654113784094
```

```
In [46]: 1 rd.predict(x[[1]])
```

```
Out[46]: array([[2.63566452]])
```

```
In [47]: 1 from sklearn.datasets import load_boston
```

```
In [49]: 1 boston = load_boston()
```

```
In [50]: 1 boston.keys()
```

```
Out[50]: dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
```

```
In [51]: 1 df = pd.DataFrame(boston.data, columns = boston.feature_names)
```

```
In [52]: 1 df.head()
```

```
Out[52]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

In [53]: 1 df.columns

Out[53]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='object')

In [54]: 1 boston.target

Out[54]: array([24. , 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1, 16.5, 18.9, 15. ,
18.9, 21.7, 20.4, 18.2, 19.9, 23.1, 17.5, 20.2, 18.2, 13.6, 19.6,
15.2, 14.5, 15.6, 13.9, 16.6, 14.8, 18.4, 21. , 12.7, 14.5, 13.2,
13.1, 13.5, 18.9, 20. , 21. , 24.7, 30.8, 34.9, 26.6, 25.3, 24.7,
21.2, 19.3, 20. , 16.6, 14.4, 19.4, 19.7, 20.5, 25. , 23.4, 18.9,
35.4, 24.7, 31.6, 23.3, 19.6, 18.7, 16. , 22.2, 25. , 33. , 23.5,
19.4, 22. , 17.4, 20.9, 24.2, 21.7, 22.8, 23.4, 24.1, 21.4, 20. ,
20.8, 21.2, 20.3, 28. , 23.9, 24.8, 22.9, 23.9, 26.6, 22.5, 22.2,
23.6, 28.7, 22.6, 22. , 22.9, 25. , 20.6, 28.4, 21.4, 38.7, 43.8,
33.2, 27.5, 26.5, 18.6, 19.3, 20.1, 19.5, 19.5, 20.4, 19.8, 19.4,
21.7, 22.8, 18.8, 18.7, 18.5, 18.3, 21.2, 19.2, 20.4, 19.3, 22. ,
20.3, 20.5, 17.3, 18.8, 21.4, 15.7, 16.2, 18. , 14.3, 19.2, 19.6,
23. , 18.4, 15.6, 18.1, 17.4, 17.1, 13.3, 17.8, 14. , 14.4, 13.4,
15.6, 11.8, 13.8, 15.6, 14.6, 17.8, 15.4, 21.5, 19.6, 15.3, 19.4,
17. , 15.6, 13.1, 41.3, 24.3, 23.3, 27. , 50. , 50. , 50. , 22.7,
25. , 50. , 23.8, 23.8, 22.3, 17.4, 19.1, 23.1, 23.6, 22.6, 29.4,
23.2, 24.6, 29.9, 37.2, 39.8, 36.2, 37.9, 32.5, 26.4, 29.6, 50. ,
32. , 29.8, 34.9, 37. , 30.5, 36.4, 31.1, 29.1, 50. , 33.3, 30.3,
34.6, 34.9, 32.9, 24.1, 42.3, 48.5, 50. , 22.6, 24.4, 22.5, 24.4,
20. , 21.7, 19.3, 22.4, 28.1, 23.7, 25. , 23.3, 28.7, 21.5, 23. ,
26.7, 21.7, 27.5, 30.1, 44.8, 50. , 37.6, 31.6, 46.7, 31.5, 24.3,
31.7, 41.7, 48.3, 29. , 24. , 25.1, 31.5, 23.7, 23.3, 22. , 20.1,
22.2, 23.7, 17.6, 18.5, 24.3, 20.5, 24.5, 26.2, 24.4, 24.8, 29.6,
42.8, 21.9, 20.9, 44. , 50. , 36. , 30.1, 33.8, 43.1, 48.8, 31. ,
36.5, 22.8, 30.7, 50. , 43.5, 20.7, 21.1, 25.2, 24.4, 35.2, 32.4,
32. , 33.2, 33.1, 29.1, 35.1, 45.4, 35.4, 46. , 50. , 32.2, 22. ,
20.1, 23.2, 22.3, 24.8, 28.5, 37.3, 27.9, 23.9, 21.7, 28.6, 27.1,
20.3, 22.5, 29. , 24.8, 22. , 26.4, 33.1, 36.1, 28.4, 33.4, 28.2,
22.8, 20.3, 16.1, 22.1, 19.4, 21.6, 23.8, 16.2, 17.8, 19.8, 23.1,
21. , 23.8, 23.1, 20.4, 18.5, 25. , 24.6, 23. , 22.2, 19.3, 22.6,
19.8, 17.1, 19.4, 22.2, 20.7, 21.1, 19.5, 18.5, 20.6, 19. , 18.7,
32.7, 16.5, 23.9, 31.2, 17.5, 17.2, 23.1, 24.5, 26.6, 22.9, 24.1,
18.6, 30.1, 18.2, 20.6, 17.8, 21.7, 22.7, 22.6, 25. , 19.9, 20.8,
16.8, 21.9, 27.5, 21.9, 23.1, 50. , 50. , 50. , 50. , 50. , 13.8,
13.8, 15. , 13.9, 13.3, 13.1, 10.2, 10.4, 10.9, 11.3, 12.3, 8.8,
7.2, 10.5, 7.4, 10.2, 11.5, 15.1, 23.2, 9.7, 13.8, 12.7, 13.1,
12.5, 8.5, 5. , 6.3, 5.6, 7.2, 12.1, 8.3, 8.5, 5. , 11.9,
27.9, 17.2, 27.5, 15. , 17.2, 17.9, 16.3, 7. , 7.2, 7.5, 10.4,
8.8, 8.4, 16.7, 14.2, 20.8, 13.4, 11.7, 8.3, 10.2, 10.9, 11. ,
9.5, 14.5, 14.1, 16.1, 14.3, 11.7, 13.4, 9.6, 8.7, 8.4, 12.8,
10.5, 17.1, 18.4, 15.4, 10.8, 11.8, 14.9, 12.6, 14.1, 13. , 13.4,
15.2, 16.1, 17.8, 14.9, 14.1, 12.7, 13.5, 14.9, 20. , 16.4, 17.7,
19.5, 20.2, 21.4, 19.9, 19. , 19.1, 19.1, 20.1, 19.9, 19.6, 23.2,
29.8, 13.8, 13.3, 16.7, 12. , 14.6, 21.4, 23. , 23.7, 25. , 21.8,
20.6, 21.2, 19.1, 20.6, 15.2, 7. , 8.1, 13.6, 20.1, 21.8, 24.5,
23.1, 19.7, 18.3, 21.2, 17.5, 16.8, 22.4, 20.6, 23.9, 22. , 11.9])

In [55]: 1 df['LandPrice']=boston.target

In [56]: 1 df.head()

Out[56]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

In [57]: 1 df.isnull().sum()

Out[57]:

CRIM	0
ZN	0
INDUS	0
CHAS	0
NOX	0
RM	0
AGE	0
DIS	0
RAD	0
TAX	0
PTRATIO	0
B	0
LSTAT	0
LandPrice	0
dtype:	int64

In [58]:

```
1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column        Non-Null Count  Dtype  
---  -
 0   CRIM          506 non-null    float64
 1   ZN            506 non-null    float64
 2   INDUS         506 non-null    float64
 3   CHAS          506 non-null    float64
 4   NOX           506 non-null    float64
 5   RM            506 non-null    float64
 6   AGE           506 non-null    float64
 7   DIS           506 non-null    float64
 8   RAD           506 non-null    float64
 9   TAX           506 non-null    float64
10  PTRATIO       506 non-null    float64
11  B             506 non-null    float64
12  LSTAT         506 non-null    float64
13  LandPrice     506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB
```

In [59]:

```
1 X = df.iloc[:,0:-1]
```

In [60]:

```
1 y=df['LandPrice']
```

In [61]:

```
1 from sklearn.linear_model import LinearRegression
2 linear = LinearRegression()
```

In [62]:

```
1 linear.fit(X,y)
```

Out[62]: LinearRegression()

In [63]:

```
1 from sklearn.model_selection import train_test_split
```

In [64]:

```
1 X_train,X_test,y_train,y_test = train_test_split(X,y,train_size = 0.75)
```

In [65]:

```
1 linear.score(X_train,y_train)
```

Out[65]: 0.7636799441572235

In [66]:

```
1 linear.score(X_test,y_test)
```

Out[66]: 0.6604205887452774

In [67]:

```
1 from sklearn.linear_model import Ridge
```

```
In [78]: 1 rd = Ridge(alpha = 0.1)
         2
```

```
In [79]: 1 rd.fit(X_train,y_train)
```

```
Out[79]: Ridge(alpha=0.1)
```

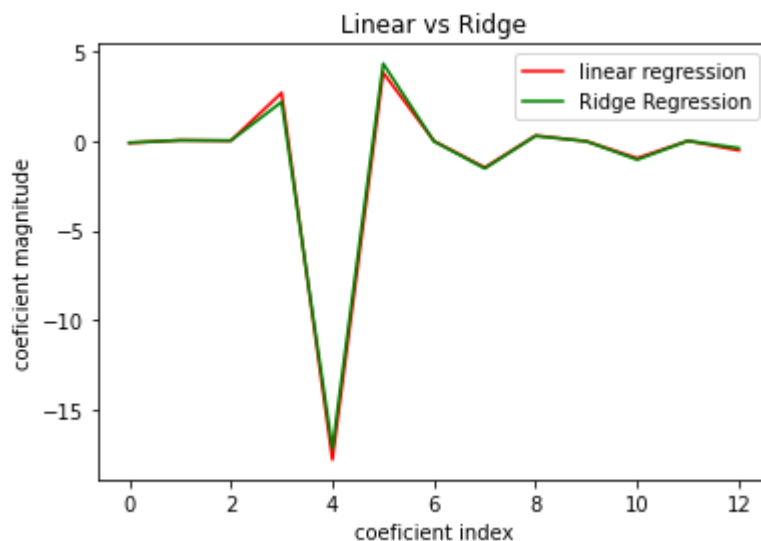
```
In [80]: 1 rd.score(X_train,y_train)*100
```

```
Out[80]: 77.00211293470278
```

```
In [81]: 1 rd.score(X_test,y_test)
```

```
Out[81]: 0.6071919674027766
```

```
In [82]: 1 plt.plot(linear.coef_,c='r',label="linear regression")
         2 plt.plot(rd.coef_,c='g',label = 'Ridge Regression')
         3 plt.title('Linear vs Ridge')
         4 plt.xlabel('coefficient index')
         5 plt.ylabel('coefficient magnitude')
         6 plt.legend()
         7 plt.show()
```



```
In [90]: 1 from sklearn.linear_model import Lasso
         2 la = Lasso(alpha =100)
```

```
In [91]: 1 la.fit(X_train,y_train)
```

```
Out[91]: Lasso(alpha=100)
```


```
In [92]: 1 la.score(X_train,y_train)
```

```
Out[92]: 0.2879575494019463
```

```
In [93]: 1 la.score(X_test,y_test)
```

```
Out[93]: 0.00949562533326187
```

```
In [94]: 1 plt.plot(linear.coef_,c='r',label="linear regression")
2 plt.plot(rd.coef_,c='g',label = 'Ridge Regression')
3 plt.plot(la.coef_,c='b',label = 'Lasso Regression')
4 plt.title('Linear vs Ridge vs Lasso')
5 plt.xlabel('coefficient index')
6 plt.ylabel('coefficient magnitude')
7 plt.legend()
8 plt.show()
```



```
In [ ]: 1
```