

Logistic Regression

```
In [1]: 1 # import data set  
        2 import pandas as pd
```

```
In [2]: 1 from sklearn import datasets
```



```

computed for each image,\n
d 3 is Mean Radius, field\n
s.\n\n
- class:\n
DBC-Benign\n\n
:Summary Statistics:\n\n
=====
===== \n
radius (mean):
6.981 28.11\n texture (mean): 9.71 39.28\n per
imeter (mean): 43.79 188.5\n area (mean):
143.5 2501.0\n smoothness (mean): 0.053 0.163\n co
mpactness (mean): 0.019 0.345\n concavity (mean):
0.0 0.427\n concave points (mean): 0.0 0.201\n sym
metry (mean): 0.106 0.304\n fractal dimension (mea
n): 0.05 0.097\n radius (standard error): 0.112
2.873\n texture (standard error): 0.36 4.885\n perimeter
(standard error): 0.757 21.98\n area (standard error):
6.802 542.2\n smoothness (standard error): 0.002 0.031\n com
pactness (standard error): 0.002 0.135\n concavity (standard erro
r): 0.0 0.396\n concave points (standard error): 0.0
0.053\n symmetry (standard error): 0.008 0.079\n fractal di
mension (standard error): 0.001 0.03\n radius (worst):
7.93 36.04\n texture (worst): 12.02 49.54\n per
imeter (worst): 50.41 251.2\n area (worst):
185.2 4254.0\n smoothness (worst): 0.071 0.223\n co
mpactness (worst): 0.027 1.058\n concavity (worst):
0.0 1.252\n concave points (worst): 0.0 0.291\n sym
metry (worst): 0.156 0.664\n fractal dimension (wors
t): 0.055 0.208\n
=====
===== \n\n
:Missing Attribute Values: None\n\n
:Class Distribution: 212
- Malignant, 357 - Benign\n\n
:Creator: Dr. William H. Wolberg, W. Nick S
treet, Olvi L. Mangasarian\n\n
:Donor: Nick Street\n\n
>Date: November,
1995\n\n
This is a copy of UCI ML Breast Cancer Wisconsin (Diagnostic) dataset
s.\n\n
https://goo.gl/U2Uwz2\n\n
Features are computed from a digitized image of
a fine needle\naspirate (FNA) of a breast mass. They describe\ncharacteristi
cs of the cell nuclei present in the image.\n\n
Separating plane described abo
ve was obtained using\nMultisurface Method-Tree (MSM-T) [K. P. Bennett, "Deci
sion Tree\nConstruction Via Linear Programming." Proceedings of the 4th\nMidw
est Artificial Intelligence and Cognitive Science Society,\npp. 97-101, 199
2], a classification method which uses linear\nprogramming to construct a dec
ision tree. Relevant features\nwere selected using an exhaustive search in t
he space of 1-4\nfeatures and 1-3 separating planes.\n\n
The actual linear pro
gram used to obtain the separating plane\nin the 3-dimensional space is that
described in:\n[K. P. Bennett and O. L. Mangasarian: "Robust Linear\nProgramm
ing Discrimination of Two Linearly Inseparable Sets",\nOptimization Methods a
nd Software 1, 1992, 23-34].\n\n
This database is also available through the U
W CS ftp server:\n\n
ftp ftp.cs.wisc.edu\ncd math-prog/cpo-dataset/machine-lea
rn/WDBC/\n\n
.. topic:: References\n\n
- W.N. Street, W.H. Wolberg and O.L.
Mangasarian. Nuclear feature extraction \n
for breast tumor diagnosis. IS
&T/SPIE 1993 International Symposium on \n
Electronic Imaging: Science an
d Technology, volume 1905, pages 861-870,\n
San Jose, CA, 1993.\n
- O.
L. Mangasarian, W.N. Street and W.H. Wolberg. Breast cancer diagnosis and \n
prognosis via linear programming. Operations Research, 43(4), pages 570-577,
\n
July-August 1995.\n
- W.H. Wolberg, W.N. Street, and O.L. Mangasaria
n. Machine learning techniques\n
to diagnose breast cancer from fine-need
le aspirates. Cancer Letters 77 (1994) \n
163-171.',
'feature_names': array(['mean radius', 'mean texture', 'mean perimeter', 'me
an area',
'mean smoothness', 'mean compactness', 'mean concavity',

```

```
'mean concave points', 'mean symmetry', 'mean fractal dimension',
'radius error', 'texture error', 'perimeter error', 'area error',
'smoothness error', 'compactness error', 'concavity error',
'concave points error', 'symmetry error',
'fractal dimension error', 'worst radius', 'worst texture',
'worst perimeter', 'worst area', 'worst smoothness',
'worst compactness', 'worst concavity', 'worst concave points',
'worst symmetry', 'worst fractal dimension'], dtype='<U23'),
'filename': 'C:\\Users\\Alekhya\\Anaconda3\\lib\\site-packages\\sklearn\\dat
assets\\data\\breast_cancer.csv'}
```

```
In [8]: 1 # selecting features and target
2 input_data = pd.DataFrame(cancer["data"], columns=['mean radius', 'mean texture',
3           'mean smoothness', 'mean compactness', 'mean concavity',
4           'mean concave points', 'mean symmetry', 'mean fractal dimension',
5           'radius error', 'texture error', 'perimeter error', 'area error',
6           'smoothness error', 'compactness error', 'concavity error',
7           'concave points error', 'symmetry error',
8           'fractal dimension error', 'worst radius', 'worst texture',
9           'worst perimeter', 'worst area', 'worst smoothness',
10          'worst compactness', 'worst concavity', 'worst concave points',
11          'worst symmetry', 'worst fractal dimension'])
12 input_data.head()
```

Out[8]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	d
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	

5 rows × 30 columns

```
In [9]: 1 input_data.shape
```

Out[9]: (569, 30)

```
In [11]: 1 output_data = pd.DataFrame(cancer["target"],columns=["target"])
        2 output_data.head()
```

Out[11]:

	target
0	0
1	0
2	0
3	0
4	0

```
In [12]: 1 output_data.shape
```

Out[12]: (569, 1)

```
In [13]: 1 input_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 30 columns):
mean radius           569 non-null float64
mean texture          569 non-null float64
mean perimeter        569 non-null float64
mean area             569 non-null float64
mean smoothness       569 non-null float64
mean compactness      569 non-null float64
mean concavity         569 non-null float64
mean concave points   569 non-null float64
mean symmetry         569 non-null float64
mean fractal dimension 569 non-null float64
radius error          569 non-null float64
texture error          569 non-null float64
perimeter error        569 non-null float64
area error             569 non-null float64
smoothness error       569 non-null float64
compactness error      569 non-null float64
concavity error        569 non-null float64
concave points error   569 non-null float64
symmetry error         569 non-null float64
fractal dimension error 569 non-null float64
worst radius           569 non-null float64
worst texture          569 non-null float64
worst perimeter        569 non-null float64
worst area             569 non-null float64
worst smoothness       569 non-null float64
worst compactness      569 non-null float64
worst concavity         569 non-null float64
worst concave points   569 non-null float64
worst symmetry         569 non-null float64
worst fractal dimension 569 non-null float64
dtypes: float64(30)
memory usage: 133.4 KB
```

```
In [14]: 1 output_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 569 entries, 0 to 568  
Data columns (total 1 columns):  
target      569 non-null int32  
dtypes: int32(1)  
memory usage: 2.3 KB
```

```
In [15]: 1 # applying preprocessing technique  
        2 input_data.isnull().sum()
```

```
Out[15]: mean radius      0  
mean texture      0  
mean perimeter    0  
mean area         0  
mean smoothness   0  
mean compactness  0  
mean concavity    0  
mean concave points 0  
mean symmetry     0  
mean fractal dimension 0  
radius error      0  
texture error     0  
perimeter error   0  
area error        0  
smoothness error  0  
compactness error 0  
concavity error   0  
concave points error 0  
symmetry error    0  
fractal dimension error 0  
worst radius      0  
worst texture     0  
worst perimeter   0  
worst area        0  
worst smoothness  0  
worst compactness 0  
worst concavity   0  
worst concave points 0  
worst symmetry    0  
worst fractal dimension 0  
dtype: int64
```

```
In [16]: 1 output_data.isnull().sum()
```

```
Out[16]: target      0  
dtype: int64
```

```
In [17]: 1 output_data["target"].value_counts()
```

```
Out[17]: 1    357  
        0    212  
        Name: target, dtype: int64
```

```
In [18]: 1 # seperating data for training and testing
          2 from sklearn.model_selection import train_test_split
```

```
In [20]: x_train,x_test,y_train,y_test = train_test_split(input_data,output_data,test
              random_state=2)
```

```
In [22]: 1 # select the model
          2 from sklearn.linear_model import LogisticRegression
```

```
In [23]: 1 log = LogisticRegression()
```

```
In [24]: 1 log.fit(x_train,y_train)
```

```
C:\Users\Alekhya\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
```

FutureWarning)

```
C:\Users\Alekhya\Anaconda3\lib\site-packages\sklearn\utils\validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

[illegible]

```
In [26]: 1 # predict the values for testing
          2 pred = log.predict(x_test)
          3 pred
```

```
Out[26]: array([1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1,
                1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1,
                0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1,
                0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1,
                1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1,
                0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0,
                1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0,
                1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1,
                0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1])
```

```
In [33]: 1 # evaluating the model
          2 from sklearn.metrics import accuracy_score, classification_report, confusion_m
```

```
In [34]: 1 accuracy_score(y_test, pred)
```

Out[34]: 0.9414893617021277

```
In [35]: 1 print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.93	0.92	0.92	73
1	0.95	0.96	0.95	115
micro avg	0.94	0.94	0.94	188
macro avg	0.94	0.94	0.94	188
weighted avg	0.94	0.94	0.94	188

```
In [36]: 1 confusion_matrix(y_test,pred)
```

```
Out[36]: array([[ 67,   6],
                [  5, 110]], dtype=int64)
```

Support Vector Machine

```
In [81]: 1 data = pd.read_csv("https://raw.githubusercontent.com/AP-State-Skill-Develop
```

```
In [82]: 1 data.head()
```

```
Out[82]:
```

	survived	pclass	name	sex	age	sibsp	parch	ticket	fare	cabin	embarked
0	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

```
In [83]: 1 data.shape
```

```
Out[83]: (891, 11)
```


In [84]: 1 data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
survived      891 non-null int64
pclass        891 non-null int64
name          891 non-null object
sex           891 non-null object
age           714 non-null float64
sibsp         891 non-null int64
parch         891 non-null int64
ticket        891 non-null object
fare          891 non-null float64
cabin         204 non-null object
embarked      889 non-null object
dtypes: float64(2), int64(4), object(5)
memory usage: 76.6+ KB
```

In [85]: 1 891-714

Out[85]: 177

In [86]: 1 891-204

Out[86]: 687

In [87]: 1 # preprocessing
2 data.isnull().sum()

Out[87]: survived 0
pclass 0
name 0
sex 0
age 177
sibsp 0
parch 0
ticket 0
fare 0
cabin 687
embarked 2
dtype: int64

In [88]: 1 data.isnull().sum().sum()

Out[88]: 866

In [89]: 1 data.drop("cabin",axis=1,inplace=True)

```
In [90]: 1 data.columns
```

```
Out[90]: Index(['survived', 'pclass', 'name', 'sex', 'age', 'sibsp', 'parch', 'ticket',  
               'fare', 'embarked'],  
              dtype='object')
```

```
In [91]: 1 data["age"]
```

```
Out[91]: 0      22.0  
         1      38.0  
         2      26.0  
         3      35.0  
         4      35.0  
         5       NaN  
         6      54.0  
         7       2.0  
         8      27.0  
         9      14.0  
        10       4.0  
        11      58.0  
        12      20.0  
        13      39.0  
        14      14.0  
        15      55.0  
        16       2.0  
        17       NaN  
        18      31.0  
        19       NaN  
        20      35.0  
        21      34.0  
        22      15.0  
        23      28.0  
        24       8.0  
        25      38.0  
        26       NaN  
        27      19.0  
        28       NaN  
        29       NaN  
        ...  
       861      21.0  
       862      48.0  
       863       NaN  
       864      24.0  
       865      42.0  
       866      27.0  
       867      31.0  
       868       NaN  
       869       4.0  
       870      26.0  
       871      47.0  
       872      33.0  
       873      47.0  
       874      28.0  
       875      15.0  
       876      20.0  
       877      19.0  
       878       NaN  
       879      56.0  
       880      25.0  
       881      33.0  
       882      22.0  
       883      28.0
```

```
884    25.0
885    39.0
886    27.0
887    19.0
888     NaN
889    26.0
890    32.0
Name: age, Length: 891, dtype: float64
```

```
In [92]: 1 data.age.mean()
```

```
Out[92]: 29.69911764705882
```

```
In [93]: 1 round(data["age"].mean())
```

```
Out[93]: 30
```

```
In [94]: 1 data["age"]=data["age"].fillna(round(data["age"].mean()))
```

```
In [95]: 1 data.isnull().sum()
```

```
Out[95]: survived    0
pclass             0
name               0
sex                0
age                0
sibsp              0
parch              0
ticket             0
fare               0
embarked           2
dtype: int64
```

```
In [96]: 1 data["embarked"].dtype
```

```
Out[96]: dtype('O')
```

```
In [97]: 1 data["embarked"].value_counts()
```

```
Out[97]: S    644
C    168
Q     77
Name: embarked, dtype: int64
```

```
In [98]: 1 data["embarked"]=data["embarked"].fillna("S")
```

```
In [99]: 1 data.isnull().sum()
```

```
Out[99]: survived    0  
pclass    0  
name    0  
sex    0  
age    0  
sibsp    0  
parch    0  
ticket    0  
fare    0  
embarked    0  
dtype: int64
```

```
In [100]: 1 data.drop("name",axis=1,inplace=True)
```

```
In [101]: 1 data.columns
```

```
Out[101]: Index(['survived', 'pclass', 'sex', 'age', 'sibsp', 'parch', 'ticket', 'fare',  
                'embarked'],  
               dtype='object')
```

```
In [102]: 1 data.shape
```

```
Out[102]: (891, 9)
```

```
In [103]: 1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 891 entries, 0 to 890  
Data columns (total 9 columns):  
survived    891 non-null int64  
pclass    891 non-null int64  
sex    891 non-null object  
age    891 non-null float64  
sibsp    891 non-null int64  
parch    891 non-null int64  
ticket    891 non-null object  
fare    891 non-null float64  
embarked    891 non-null object  
dtypes: float64(2), int64(4), object(3)  
memory usage: 62.7+ KB
```

In [104]:

```
1 data.head()
```

Out[104]:

	survived	pclass	sex	age	sibsp	parch	ticket	fare	embarked
0	0	3	male	22.0	1	0	A/5 21171	7.2500	S
1	1	1	female	38.0	1	0	PC 17599	71.2833	C
2	1	3	female	26.0	0	0	STON/O2. 3101282	7.9250	S
3	1	1	female	35.0	1	0	113803	53.1000	S
4	0	3	male	35.0	0	0	373450	8.0500	S

```
In [105]: 1 data["ticket"].value_counts()
```

```
Out[105]: CA. 2343          7
          347082          7
          1601           7
          347088          6
          CA 2144          6
          3101295         6
          382652          5
          S.O.C. 14879    5
          113781          4
          347077          4
          PC 17757        4
          17421           4
          349909          4
          2666            4
          19950           4
          4133            4
          W./C. 6608      4
          113760          4
          LINE            4
          345773          3
          13502           3
          363291          3
          C.A. 31921      3
          PC 17760        3
          PC 17572        3
          C.A. 34651      3
          35273           3
          110152          3
          230080          3
          239853          3
          ..
          233639          1
          349221          1
          F.C.C. 13528    1
          2687            1
          F.C.C. 13531    1
          SC/AH Basle 541 1
          19952           1
          349243          1
          2672            1
          368323          1
          65303           1
          2647            1
          PC 17601        1
          350034          1
          250648          1
          250651          1
          13509           1
          349223          1
          349222          1
          4135            1
          113786          1
          350026          1
          228414          1
```

```
112050      1
65304       1
C.A. 24580  1
367232      1
330919      1
14973       1
7267        1
Name: ticket, Length: 681, dtype: int64
```

```
In [106]: 1 data.drop("ticket",axis=1,inplace=True)
```

```
In [107]: 1 data.columns
```

```
Out[107]: Index(['survived', 'pclass', 'sex', 'age', 'sibsp', 'parch', 'fare',
                'embarked'],
                dtype='object')
```

```
In [108]: 1 from sklearn.preprocessing import LabelEncoder
```

```
In [109]: 1 lab = LabelEncoder()
```

```
In [110]: 1 data["sex"] = lab.fit_transform(data["sex"])
          2 data["sex"].head()
```

```
Out[110]: 0    1
          1    0
          2    0
          3    0
          4    1
          Name: sex, dtype: int32
```

```
In [111]: 1 data["embarked"] = lab.fit_transform(data["embarked"])
          2 data["embarked"].head()
```

```
Out[111]: 0    2
          1    0
          2    2
          3    2
          4    2
          Name: embarked, dtype: int32
```


In [112]: 1 data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 8 columns):
survived      891 non-null int64
pclass        891 non-null int64
sex           891 non-null int32
age           891 non-null float64
sibsp         891 non-null int64
parch         891 non-null int64
fare          891 non-null float64
embarked      891 non-null int32
dtypes: float64(2), int32(2), int64(4)
memory usage: 48.8 KB
```

In [113]: 1 data.head()

Out[113]:

	survived	pclass	sex	age	sibsp	parch	fare	embarked
0	0	3	1	22.0	1	0	7.2500	2
1	1	1	0	38.0	1	0	71.2833	0
2	1	3	0	26.0	0	0	7.9250	2
3	1	1	0	35.0	1	0	53.1000	2
4	0	3	1	35.0	0	0	8.0500	2

In [114]: 1 data.columns

Out[114]: Index(['survived', 'pclass', 'sex', 'age', 'sibsp', 'parch', 'fare',
'embarked'],
dtype='object')

In [115]: 1 *# selecting features and target*
2 input1 = data.drop("survived",axis=1)
3 input1.head()

Out[115]:

	pclass	sex	age	sibsp	parch	fare	embarked
0	3	1	22.0	1	0	7.2500	2
1	1	0	38.0	1	0	71.2833	0
2	3	0	26.0	0	0	7.9250	2
3	1	0	35.0	1	0	53.1000	2
4	3	1	35.0	0	0	8.0500	2

```
In [116]: 1 data.columns
```

```
Out[116]: Index(['survived', 'pclass', 'sex', 'age', 'sibsp', 'parch', 'fare',  
                'embarked'],  
                dtype='object')
```

```
In [117]: 1 output1 = data["survived"]  
          2 output1.head()
```

```
Out[117]: 0    0  
          1    1  
          2    1  
          3    1  
          4    0  
          Name: survived, dtype: int64
```

```
In [119]: 1 # seperating data for training and testing  
          2 from sklearn.model_selection import train_test_split
```

```
In [132]: 1 x_train,x_test,y_train,y_test = train_test_split(input1,output1,test_size=0.
```

In [133]:

```

1 # select the model
2 from sklearn.svm import SVC
3 help(SVC)

```

Help on class SVC in module sklearn.svm.classes:

```

class SVC(sklearn.svm.base.BaseSVC)
| SVC(C=1.0, kernel='rbf', degree=3, gamma='auto_deprecated', coef0=0.0, sh
| rinking=True, probability=False, tol=0.001, cache_size=200, class_weight=Non
| e, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=No
| ne)
|
| C-Support Vector Classification.
|
| The implementation is based on libsvm. The fit time complexity
| is more than quadratic with the number of samples which makes it hard
| to scale to dataset with more than a couple of 10000 samples.
|
| The multiclass support is handled according to a one-vs-one scheme.
|
| For details on the precise mathematical formulation of the provided
| kernel functions and how `gamma`, `coef0` and `degree` affect each
| other, see the corresponding section in the narrative documentation:
| :ref:`svm_kernels`.
|
| Read more in the :ref:`User Guide <svm_classification>`.
|
| Parameters
| -----
| C : float, optional (default=1.0)
|     Penalty parameter C of the error term.
|
| kernel : string, optional (default='rbf')
|     Specifies the kernel type to be used in the algorithm.
|     It must be one of 'linear', 'poly', 'rbf', 'sigmoid', 'precomputed' o
|
|     a callable.
|     If none is given, 'rbf' will be used. If a callable is given it is
|     used to pre-compute the kernel matrix from data matrices; that matrix
|     should be an array of shape ``(n_samples, n_samples)``.
|
| degree : int, optional (default=3)
|     Degree of the polynomial kernel function ('poly').
|     Ignored by all other kernels.
|
| gamma : float, optional (default='auto')
|     Kernel coefficient for 'rbf', 'poly' and 'sigmoid'.
|
|     Current default is 'auto' which uses 1 / n_features,
|     if ``gamma='scale'`` is passed then it uses 1 / (n_features * X.var
|
|     as value of gamma. The current default of gamma, 'auto', will change
|     to 'scale' in version 0.22. 'auto_deprecated', a deprecated version o
|
|     'auto' is used as a default indicating that no explicit value of gamm
|
| a

```

was passed.

```
coef0 : float, optional (default=0.0)
    Independent term in kernel function.
    It is only significant in 'poly' and 'sigmoid'.

shrinking : boolean, optional (default=True)
    Whether to use the shrinking heuristic.

probability : boolean, optional (default=False)
    Whether to enable probability estimates. This must be enabled prior
    to calling `fit`, and will slow down that method.

tol : float, optional (default=1e-3)
    Tolerance for stopping criterion.

cache_size : float, optional
    Specify the size of the kernel cache (in MB).

class_weight : {dict, 'balanced'}, optional
    Set the parameter C of class i to class_weight[i]*C for
    SVC. If not given, all classes are supposed to have
    weight one.
    The "balanced" mode uses the values of y to automatically adjust
    weights inversely proportional to class frequencies in the input data
    as n_samples / (n_classes * np.bincount(y))

verbose : bool, default: False
    Enable verbose output. Note that this setting takes advantage of a
    per-process runtime setting in libsvm that, if enabled, may not work
    properly in a multithreaded context.

max_iter : int, optional (default=-1)
    Hard limit on iterations within solver, or -1 for no limit.

decision_function_shape : 'ovo', 'ovr', default='ovr'
    Whether to return a one-vs-rest ('ovr') decision function of shape
    (n_samples, n_classes) as all other classifiers, or the original
    one-vs-one ('ovo') decision function of libsvm which has shape
    (n_samples, n_classes * (n_classes - 1) / 2). However, one-vs-one
    ('ovo') is always used as multi-class strategy.

.. versionchanged:: 0.19
    decision_function_shape is 'ovr' by default.

.. versionadded:: 0.17
    *decision_function_shape='ovr'* is recommended.

.. versionchanged:: 0.17
    Deprecate *decision_function_shape='ovo' and None*.

random_state : int, RandomState instance or None, optional (default=None)
    The seed of the pseudo random number generator used when shuffling
    the data for probability estimates. If int, random_state is the
    seed used by the random number generator; If RandomState instance,
    random_state is the random number generator; If None, the random
    number generator is the RandomState instance used by `np.random`.
```

Attributes

`support_` : array-like, shape = `[n_SV]`
Indices of support vectors.

`support_vectors_` : array-like, shape = `[n_SV, n_features]`
Support vectors.

`n_support_` : array-like, dtype=int32, shape = `[n_class]`
Number of support vectors for each class.

`dual_coef_` : array, shape = `[n_class-1, n_SV]`
Coefficients of the support vector in the decision function.
For multiclass, coefficient for all 1-vs-1 classifiers.
The layout of the coefficients in the multiclass case is somewhat non-trivial. See the section about multi-class classification in the SVM section of the User Guide for details.

`coef_` : array, shape = `[n_class * (n_class-1) / 2, n_features]`
Weights assigned to the features (coefficients in the primal problem). This is only available in the case of a linear kernel.

``coef_`` is a readonly property derived from ``dual_coef_`` and ``support_vectors_``.

`intercept_` : array, shape = `[n_class * (n_class-1) / 2]`
Constants in decision function.

`fit_status_` : int
0 if correctly fitted, 1 otherwise (will raise warning)

`probA_` : array, shape = `[n_class * (n_class-1) / 2]`

`probB_` : array, shape = `[n_class * (n_class-1) / 2]`

If `probability=True`, the parameters learned in Platt scaling to produce probability estimates from decision values. If `probability=False`, an empty array. Platt scaling uses the logistic function

``1 / (1 + exp(decision_value * probA_ + probB_))``

where ``probA_`` and ``probB_`` are learned from the dataset `[2]_`. For

more information on the multiclass case and training procedure see section 8 of `[1]_`.

Examples

```
>>> import numpy as np
>>> X = np.array([[-1, -1], [-2, -1], [1, 1], [2, 1]])
>>> y = np.array([1, 1, 2, 2])
>>> from sklearn.svm import SVC
>>> clf = SVC(gamma='auto')
>>> clf.fit(X, y) #doctest: +NORMALIZE_WHITESPACE
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
>>> print(clf.predict([[-0.8, -1]]))
```

[1]

See also

SVR

Support Vector Machine for Regression implemented using libsvm.

LinearSVC

Scalable Linear Support Vector Machine for classification implemented using liblinear. Check the See also section of LinearSVC for more comparison element.

References

.. [1] `LIBSVM: A Library for Support Vector Machines

<<http://www.csie.ntu.edu.tw/~cjlin/papers/libsvm.pdf>>`_

.. [2] `Platt, John (1999). "Probabilistic outputs for support vector machines and comparison to regularized likelihood methods."

<<http://citeseer.ist.psu.edu/viewdoc/summary?doi=10.1.1.41.1639>>`_

Method resolution order:

SVC

sklearn.svm.base.BaseSVC

abc.NewBase

sklearn.svm.base.BaseLibSVM

abc.NewBase

sklearn.base.BaseEstimator

sklearn.base.ClassifierMixin

builtins.object

Methods defined here:

```
__init__(self, C=1.0, kernel='rbf', degree=3, gamma='auto_deprecated', coef0=0.0, shrinking=True, probability=False, tol=0.001, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=None)
```

Initialize self. See help(type(self)) for accurate signature.

Data and other attributes defined here:

```
__abstractmethods__ = frozenset()
```

Methods inherited from sklearn.svm.base.BaseSVC:

```
decision_function(self, X)
```

Evaluates the decision function for the samples in X.

Parameters

X : array-like, shape (n_samples, n_features)

Returns

X : array-like, shape (n_samples, n_classes * (n_classes-1) / 2)

Returns the decision function of the sample for each class in the model.
If `decision_function_shape='ovr'`, the shape is `(n_samples, n_classes)`.

Notes

If `decision_function_shape='ovo'`, the function values are proportional to the distance of the samples `X` to the separating hyperplane. If the exact distances are required, divide the function values by the norm of the weight vector (`coef_`). See also `this question <https://stats.stackexchange.com/questions/14876/interpreting-distance-from-hyperplane-in-svm>` for further details.

`predict(self, X)`

Perform classification on samples in `X`.

For an one-class model, `+1` or `-1` is returned.

Parameters

`X` : {array-like, sparse matrix}, shape `(n_samples, n_features)`
For `kernel="precomputed"`, the expected shape of `X` is `[n_samples_test, n_samples_train]`

Returns

`y_pred` : array, shape `(n_samples,)`
Class labels for samples in `X`.

Data descriptors inherited from `sklearn.svm.base.BaseSVC`:

`predict_log_proba`

Compute log probabilities of possible outcomes for samples in `X`.

The model need to have probability information computed at training time: fit with attribute `'probability'` set to `True`.

Parameters

`X` : array-like, shape `(n_samples, n_features)`
For `kernel="precomputed"`, the expected shape of `X` is `[n_samples_test, n_samples_train]`

Returns

`T` : array-like, shape `(n_samples, n_classes)`
Returns the log-probabilities of the sample for each class in the model. The columns correspond to the classes in sorted order, as they appear in the attribute `'classes_'`.

Notes

The probability model is created using cross validation, so

the results can be slightly different than those obtained by predict. Also, it will produce meaningless results on very small datasets.

predict_proba

Compute probabilities of possible outcomes for samples in X.

The model need to have probability information computed at training time: fit with attribute `probability` set to True.

Parameters

X : array-like, shape (n_samples, n_features)
For kernel="precomputed", the expected shape of X is [n_samples_test, n_samples_train]

Returns

T : array-like, shape (n_samples, n_classes)
Returns the probability of the sample for each class in the model. The columns correspond to the classes in sorted order, as they appear in the attribute `classes_`.

Notes

The probability model is created using cross validation, so the results can be slightly different than those obtained by predict. Also, it will produce meaningless results on very small datasets.

Methods inherited from sklearn.svm.base.BaseLibSVM:

fit(self, X, y, sample_weight=None)

Fit the SVM model according to the given training data.

Parameters

X : {array-like, sparse matrix}, shape (n_samples, n_features)
Training vectors, where n_samples is the number of samples and n_features is the number of features.
For kernel="precomputed", the expected shape of X is (n_samples, n_samples).

y : array-like, shape (n_samples,)
Target values (class labels in classification, real numbers in regression)

sample_weight : array-like, shape (n_samples,)
Per-sample weights. Rescale C per sample. Higher weights force the classifier to put more emphasis on these points.

Returns

self : object

Notes

If X and y are not C-ordered and contiguous arrays of np.float64 and X is not a scipy.sparse.csr_matrix, X and/or y may be copied.

If X is a dense array, then the other methods will not support sparse matrices as input.

Data descriptors inherited from sklearn.svm.base.BaseLibSVM:

coef_

Methods inherited from sklearn.base.BaseEstimator:

__getstate__(self)

__repr__(self)
Return repr(self).

__setstate__(self, state)

get_params(self, deep=True)
Get parameters for this estimator.

Parameters

deep : boolean, optional
If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns

params : mapping of string to any
Parameter names mapped to their values.

set_params(self, **params)
Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as pipelines). The latter have parameters of the form ``<component>__<parameter>`` so that it's possible to update each component of a nested object.

Returns

self

Data descriptors inherited from sklearn.base.BaseEstimator:

__dict__
dictionary for instance variables (if defined)

__weakref__
list of weak references to the object (if defined)

 Methods inherited from sklearn.base.ClassifierMixin:

`score(self, X, y, sample_weight=None)`

Returns the mean accuracy on the given test data and labels.

In multi-label classification, this is the subset accuracy which is a harsh metric since you require for each sample that each label set be correctly predicted.

Parameters

`X` : array-like, shape = (n_samples, n_features)

Test samples.

`y` : array-like, shape = (n_samples) or (n_samples, n_outputs)

True labels for X.

`sample_weight` : array-like, shape = [n_samples], optional

Sample weights.

Returns

`score` : float

Mean accuracy of `self.predict(X)` wrt. `y`.

In [134]: 1 `sv = SVC(kernel="linear")`

In [135]: 1 `# train the model`
 2 `sv.fit(x_train,y_train)`

Out[135]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0, decision_function_shape='ovr', degree=3, gamma='auto_deprecated', kernel='linear', max_iter=-1, probability=False, random_state=None, shrinking=True, tol=0.001, verbose=False)

```
In [136]: 1 #predicting
          2 p = sv.predict(x_test)
          3 p
```

```
Out[136]: array([0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1,
                0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0,
                0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0,
                1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0,
                1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0,
                0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
                1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0,
                1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0,
                0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0,
                0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0,
                0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0,
                0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
                0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0,
                1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1,
                0, 1, 0, 1], dtype=int64)
```

```
In [137]: 1 from sklearn.metrics import accuracy_score, confusion_matrix, classification_r
```

```
In [138]: 1 accuracy_score(y_test, p)
```

```
Out[138]: 0.7756410256410257
```

```
In [139]: 1 confusion_matrix(y_test, p)
```

```
Out[139]: array([[159, 24],
                 [ 46, 83]], dtype=int64)
```

```
In [140]: 1 print(classification_report(y_test, p))
```

	precision	recall	f1-score	support
0	0.78	0.87	0.82	183
1	0.78	0.64	0.70	129
micro avg	0.78	0.78	0.78	312
macro avg	0.78	0.76	0.76	312
weighted avg	0.78	0.78	0.77	312

```
In [ ]: 1
```

```
In [ ]: 1
```